

Unified Applications of Generative and Transformer-based Models Across Modalities

Abstract

This project explores how **generative and Transformer-based AI models** can solve complex problems across diverse data types. We focus on four key areas: **musical source separation** (splitting songs into individual instruments), **real-time video matting** (separating foreground from background in live video), **protein structure prediction** (determining protein shapes from their sequences), and **few-shot multi-modal classification** (teaching AI to learn new concepts with very few examples across images and text).

We use leading models like **Demucs, MODNet, AlphaFold2, CLIP, and Flamingo**, alongside their standard datasets. Our goal is to see how well these models **transfer** their learning to new domains, **generalize** to limited or new data, and what their **limitations** are in real-world scenarios. This research aims to offer practical insights into the capabilities and future of these powerful AI systems.

1. Introduction

Transformers and **generative models** have revolutionized AI. Transformers, with their self-attention mechanism, excel at understanding complex relationships in data like text, images, and even biological sequences. Generative models are great at creating new, realistic data.

This project combines these strengths to tackle challenges in four distinct fields:

1. **Musical Source Separation:** Separating vocals, drums, etc., from a mixed song.
2. **Real-time Video Matting:** Precisely cutting out subjects from video backgrounds.
3. **Protein Structure Prediction:** Uncovering the 3D shapes of proteins.
4. **Few-shot Multi-modal Classification:** Enabling AI to learn new categories with minimal examples, blending visual and textual understanding.

Our aim is to understand how well these advanced models **generalize** to new data, **transfer knowledge** between different areas, and perform robustly in real-world, complex situations.

2. Related Work Overview

Our project builds on the success of several groundbreaking models:

- **Transformers [1]:** This core architecture uses "attention" to process data in parallel, efficiently understanding long-range dependencies in sequences. It's the foundation for many modern AI breakthroughs.
 - **Demucs [2]:** A deep learning model for music separation that processes both audio waveforms and spectrograms to accurately split mixed tracks into individual instruments.
 - **MODNet [10]:** A real-time network for video matting, designed to quickly and precisely separate foreground objects from backgrounds, even for complex details like hair.
 - **AlphaFold2 [3]:** A revolutionary model that accurately predicts the 3D structure of proteins from their amino acid sequences, using a sophisticated "Evoformer" architecture.
 - **CLIP [4] and Flamingo [5]:**
 - **CLIP** learns strong connections between images and text from vast datasets, allowing it to classify objects without specific training examples (zero-shot learning).
 - **Flamingo** extends CLIP by integrating large language models with visual inputs, enabling it to quickly learn new tasks from just a few examples (few-shot learning) by processing both images and text interchangeably.
-

3. Data Collection & Characteristics

To evaluate these models, we use established datasets:

- **MUSDB18 [6]:** For music separation, this dataset provides 150 full-length songs with individually separated instrument tracks.
 - **YouTube-VOS [7]:** For video matting, it offers thousands of video clips with precise object masks, crucial for training and testing real-time matting.
 - **ProteinNet [8]:** For protein prediction, it provides standardized protein sequences and their corresponding 3D structures.
 - **LAION-400M [9]:** A massive dataset of 400 million image-text pairs, used for training large multi-modal models like CLIP and Flamingo, enabling them to understand the relationship between vision and language.
-

4. Initial Findings (Base Model Implementation)

We started by setting up and testing **Demucs** for musical source separation.

- **Implementation:** We used a PyTorch implementation of Demucs, processing raw audio waveforms. We prepared a subset of the MUSDB18 dataset, handling audio pre-processing like resampling and normalization.
 - **Observations:** Even with a pre-trained model, the separation quality for vocals and drums was very promising. Our audio pre-processing system proved robust. The modular setup suggests good scalability for future experiments. However, we noted that running Demucs is computationally intensive, highlighting the need for powerful hardware for full training. We also observed minor artifacts in the separated "other instruments" track, which suggests areas for future refinement.
-

5. Formal Description of Key Algorithms

Our project relies on these advanced algorithms:

- **Transformers [1]:** Their core is the "attention mechanism," which allows the model to weigh the importance of different parts of the input, enabling it to understand long-range dependencies and process data in parallel. Multi-head attention and positional encodings further enhance their capabilities.
 - **Generative Models:** These models learn the underlying data distribution to create new, realistic data. Examples include VAEs (Variational Autoencoders) and GANs (Generative Adversarial Networks), and more recently, Diffusion Models, which iteratively refine noisy data into coherent samples.
 - **Demucs [2]:** This model uses a U-Net-like architecture with 1D and 2D convolutions to effectively separate audio sources by processing both waveform and spectrogram data.
 - **MODNet [10]:** A three-branch network that accurately performs real-time portrait matting, even without explicit trimaps.
 - **AlphaFold2 [3]:** Revolutionized protein structure prediction using its "Evoformer" block, a Transformer-like architecture that refines 3D coordinates by integrating evolutionary and structural data.
 - **CLIP [4] & Flamingo [5]:**
 - **CLIP** learns powerful visual representations by aligning images with text descriptions using a contrastive learning approach, allowing for zero-shot classification.
 - **Flamingo** builds on CLIP, integrating large language models with visual encoders to enable rapid few-shot learning by processing interleaved image and text inputs.
-

6. Unfinished Work & Future Plans

Unfinished Sections

We still need to:

- Fully evaluate MODNet for video matting, AlphaFold2 for protein folding, and CLIP/Flamingo for few-shot multi-modal classification.
- Fine-tune and optimize each model for peak performance.
- Conduct in-depth analysis of how these models work internally using visualizations.
- Perform rigorous statistical validation of our results.
- Write a comprehensive final report with detailed interpretations.

Anticipated Difficulties

Major challenges include:

- **Computational resources:** Training large models requires significant GPU power and can be costly.
- **Data complexity:** Managing and pre-processing diverse, large-scale datasets is challenging.
- **Model complexity:** Implementing and interpreting these advanced "black-box" models is difficult.
- **Cross-domain analysis:** Integrating insights across very different AI domains requires deep conceptual understanding.

Plans for Final Report

Our final report will offer:

- **Comprehensive evaluation:** Both quantitative metrics and qualitative examples for each model and task.
- **Detailed analysis:** Discussing strengths, limitations, and failure modes across all modalities.
- **Enhanced visualizations:** To intuitively demonstrate model behavior.
- **Real-world applications:** Exploring the practical implications in various industries and societal impacts.
- **Future research directions:** Outlining next steps to build more robust and versatile AI systems.

References

[1] Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł., & Polosukhin, I. (2017). Attention Is All You Need. *Advances in Neural Information Processing Systems*, 30.

- [2] Défossez, A., Usunier, N., Bottou, L., & Schwenk, H. (2020). Demucs: Music Source Separation in the Waveform Domain. *International Conference on Learning Representations*.
- [3] Jumper, J., Evans, R., Pritzel, A., Green, T., Figurnov, M., Ronneberger, O., Tunyasuvunakool, K., Bates, R., Žídek, A., Potapenko, A., Bridgland, A., Meyer, A., Kohl, S. A. A., Ballard, A. J., Cowie, A., Romera-Paredes, B., Nikolov, V., Jain, R., Adler, H., ... Hassabis, D. (2021). Highly accurate protein structure prediction with AlphaFold. *Nature*, 596(7873), 583–589.
- [4] Radford, A., Kim, J. W., Hallacy, C., Ramesh, A., Goh, G., Agarwal, S., Sastry, G., Askell, S., Mishkin, P., Clark, J., Krueger, G., & Sutskever, I. (2021). Learning Transferable Visual Models From Natural Language Supervision. *Proceedings of the 38th International Conference on Machine Learning*.
- [5] Alayrac, J.-B., Donahue, J., Luc, V., Miech, A., Ritter, I., Singh, J., Simonyan, K., Ring, R., Bulat, A., Mahajan, A., Han, B., & Lenc, K. (2022). Flamingo: A Visual Language Model for Few-Shot Learning. *Advances in Neural Information Processing Systems*.
- [6] Rafii, Z., Liutkus, A., Stöter, F. R., Mimitakis, S. I. G., & Bittner, R. (2017). The MUSDB18 Dataset for Music Source Separation. *Proceedings of the 18th International Society for Music Information Retrieval Conference*.
- [7] Xu, N., Jiang, M., Zhang, B., Huang, X., Zhao, P., Zhang, J., & Liu, X. (2017). YouTube-VOS: A Large-scale Video Object Segmentation Benchmark. *IEEE International Conference on Computer Vision (ICCV)*.
- [8] Al-Qadomi, E. H., & Barash, D. (2019). ProteinNet: A Standardized Dataset for Protein Structure Machine Learning. *Journal of Chemical Information and Modeling*, 59(12), 4967–4978.
- [9] Schuhmann, C., Beaumont, R., Gordon, R., Wightman, M., Cherti, M., Coombes, T., Katta, K., Jitsev, J., & Komatsuzaki, P. (2021). LAION-400M: Open Dataset of CLIP-Filtered 400 Million Image-Text Pairs. *arXiv preprint arXiv:2111.02114*.
- [10] Ke, L., Li, B., Zhou, Y., Zhang, R., Li, S., & Li, H. (2020). MODNet: Is a Green Screen Really Necessary for Real-time Portrait Matting? *Proceedings of the 28th ACM International Conference on Multimedia*.
- [11] Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., Chanan, G., Killeen, T., Lin, Z., Singh, N., Somers, A., Yang, J., & Chintala, S. (2019). PyTorch: An Imperative Style, High-Performance Deep Learning Library. *Advances in Neural Information Processing Systems*, 32.
- [12] Kingma, D. P., & Welling, M. (2014). Auto-Encoding Variational Bayes. *International Conference on Learning Representations*.

[13] Goodfellow, I., Pouget-Abadie, J., Mirza, M., Xu, B., Warde-Farley, D., Ozair, S., Courville, A., & Bengio, Y. (2014). Generative Adversarial Nets. *Advances in Neural Information Processing Systems*, 27.

[14] Ho, J., Jain, A., & Abbeel, P. (2020). Denoising Diffusion Probabilistic Models. *Advances in Neural Information Processing Systems*, 33.